

Dr. Sophia Chen

sophia.chen@mlresearch.io | (555) 789-0123 | New York, NY | linkedin.com/in/sophiachen-phd | github.com/sophiachen-ml

PROFESSIONAL SUMMARY

Lead Machine Learning Engineer with 6+ years of expertise delivering production deep learning and predictive models. Specialized in NLP, LLM fine-tuning, RAG pipelines, and scalable feature engineering on distributed GPU clusters.

TECHNICAL SKILLS

Core: Python, R, C++, PyTorch, TensorFlow, Scikit-learn, HuggingFace, LangChain Data & Cloud: Spark, Kafka, Snowflake, GCP Vertex AI, AWS SageMaker, Docker, MLflow Techniques: Large Language Models, Transformer Architectures, Semantic Search, Time-Series Forecasting

WORK EXPERIENCE

Lead Machine Learning Scientist | Omniscient AI (2022 · Present)

- Engineered enterprise RAG pipeline utilizing embedding models and vector databases, boosting document retrieval precision by 42% for 85k enterprise users.
- Fine-tuned open-source LLMs on proprietary financial corpora, reducing GPU inference overhead by 28% while improving accuracy by 15%.
- Deployed real-time fraud detection pipeline processing \$50M in daily transactions with under 40ms inference latency.

Data Scientist | Horizon Analytics (2019 · 2022)

- Built predictive customer churn model that retained \$2.4M in annual subscription revenue across 120 enterprise accounts.
- Spearheaded A/B testing infrastructure evaluating recommendation algorithms, yielding an 18% lift in user engagement.

EDUCATION

Ph.D. in Computer Science (Machine Learning Focus) | Columbia University (2015 · 2019)

B.S. in Applied Mathematics | MIT (2011 · 2015)